Improving sample-efficiency of model-free reinforcement learning algorithms by learning latent space representations

A systematic analysis of leveraging state representation learning for more efficient model-free reinforcement learning

Master's thesis in Computer science and engineering

Marko Guberina, Betelhem Dejene Desta

Improving sample-efficiency of model-free reinforcement learning algorithms by learning latent space representations

A systematic analysis of leveraging state representation learning for more efficient model-free reinforcement learning

Marko Guberina, Betelhem Dejene Desta

UNIVERSITY OF
GOTHENBURG

CHALMERS
UNIVERSITY OF TECHNOLOGY

Improving sample-efficiency of model-free reinforcement learning algorithms by learning latent space representations
A systematic analysis of leveraging state representation
learning for more efficient model-free reinforcement learning
Marko Guberina, Betelhem Dejene Desta

Improving sample-efficiency of model-free reinforcement learning algorithms by learning latent space representations
A systematic analysis of leveraging state representation
learning for more efficient model-free reinforcement learning
Marko Guberina, Betelhem Dejene Desta
Department of Computer Science and Engineering
Chalmers University of Technology and University of Gothenburg

Abstract

Reinforcement learning struggles to solve control tasks on directly on images. Performance on identical tasks with access to the underlying states is much better. One avenue to bridge the gap between the two is to leverage unsupervised learning as a means of learning state representations from images, thereby resulting in a better conditioned reinforcement learning problem. Through investigation of related work, characteristics of successful integration of unsupervised learning and reinforcement learning are identified. We hypothesize that joint training of state representations and policies result in highest sample-efficiency if adequate regularization is provided. We further hypothesize that representations which correlate more strongly with the underlying Markov decision process result in additional sample-efficiency. These hypotheses are tested through a simple deterministic generative representation learning model trained with image reconstruction loss and additional forward and inverse auxiliary losses. While our algorithm does not reach state-of-the-art performance, its modular implementation integrated in the reinforcement learning library Tianshou enables easy use to reinforcement learning practitioners, and thus also accelerates further research. In our tests we limited ourselves to Atari environments and primarily used Rainbow as the underlying reinforcement learning algorithm.

Acknowledgements

We give special thanks to the supervisor, examiner and everyone else at Chalmers who made this work possible.

Contents

Contents

List of Figures

List of Figures

List of Tables

List of Tables

1
Introduction

1.1　What is reinforcement learning?

In computer science, reinforcement learning is the formalization of trial-and-error learning. While this is not the only legitimate interpretation of the concept, it is the most straightforward one: "trial" implies existence of an agent which observes its environment and interacts with it though its own actions. "Error" implies that the agent has a goal it tries to achieve and that it does not know how to achieve it (in the most effective manner). What it can do is take different actions and appraise them according to how closely they lead the agent toward its goal, thereby observing the quality of those actions. By repeatedly exploring the effects of various sequences of actions, the agent can find, i.e. learn, the sequence of actions which lead to its goal.

Here, it is important to discuss what a goal is. To formalize the process outlined above, one needs to describe it in purely mathematical terms. Thus, among other things, the goal needs to be described numerically. To achieve that, the notion of a reward function is used: it maps every state of the environment to a number which denotes its value called the reward. The state of the environment to which the highest reward is ascribed is then the goal. A more general description of the goal of reinforcement learning is to maximize the sum of rewards over time. The formalization of the entire process will be carried out later in 2, while here only the most important concepts will be outlined.

Due to its generality, reinforcement learning is studied in many different disciplines: control theory, game theory, information theory, simulation-based optimization, multi-agent systems etc.. Of these, control theory is of particular importance because it often enables clear analysis of various reinforcement learning algorithms. This foremost concerns the usage of dynamic programming which provides a basis for a large class of reinforcement learning algorithms. Reinforcement learning is also considered to be one of the pillars of modern data-driven machine learning.

In the context of machine learning, reinforcement learning can be view as a combination of supervised and unsupervised learning: the "trial" portion of the trial-and-error learning can be interpreted as unsupervised or as self-supervised learning because in it the agent collects its own dataset without any explicit labels to guide its way. This process is referred to as "exploration". The dataset created by exploration is labelled by the reward function. Thus the agent can learning from "past experience" in a supervised manner. This text will introduce concepts from both control theory and machine learning which are necessary to formalize the reinforce-

ment learning objective and to develop algorithms to achieve it. It will not concern itself with other disciplines.

1.2 Why is reinforcement learning interesting?

Interest in reinforcement learning has grown tremendously over the past decade. It has been fueled by successes of deep machine learning in fields such as computer vision. The subsequent utilization of neural networks in reinforcement learning, dubbed deep reinforcement learning, led to impressive results such as achieving better-than-human performance on Atari games, in the game of go and in many others. Because large amounts of data are required for neural network training and thus for reinforcement learning algorithms which utilize them, most of these results are achieved in computer-simulated environments. [1]

While there is case to be made that reinforcement learning is a step toward artificial general intelligence, there are also more immediate applications. Due to generality of reinforcement learning and to advances in computer hardware, reinforcement learning offers a promising avenue toward solving decision and control problems which have not been solved through other, more direct methods. One of these is the problem of robotic grasping. Humans and other animals have an intuitive understanding of physics which they leverage for object manipulation. On the other hand, to program a robot to do the same, exact physics equations need to be provided so that the robot's actions may be calculated. Owing to the complexity of contact dynamics and the inability to precisely measure the points of contact, this is often impossible to do. By learning through trial and error and by leveraging the strong interpolation capabilities of neural networks, such an "intuition" may be learned. Furthermore, traditional optimization methods require rich objective functions at every iteration step, while reinforcement learning can handle "sparse" rewards — objective functions which equate to 0 at nearly all points of the domain.

1.3 Why learn from pixels?

Recent success in reinforcement learning were kick-started by Deep Q-Network (DQN) algorithm [Mni+13] which crucially, by utilizing convolutional neural network, enabled the agents to successfully learn from raw pixels. Learning from pixels is incredibly important for many practical applications, such as those in robotics where it is often impossible to get full access to the state of the environment. The state then needs to be inferred from observations such as those from cameras. Here the state refers to the underlying physical parameters of the environment: the positions and velocities of objects, the friction coefficients and so on. Observations from sensors such as cameras do not explicitly provide such information. However, since humans and animals are able to utilize such observations to achieve their goals, we know that they implicitly hold enough information about the true state of the world for successful goal completion.

---

[1]Simulated environments run as fast as the computers they run on, which enables generating thousands of trials in seconds.

The problem is that pixel-based observations are much higher-dimensional than the actual states. This makes the learning problem dramatically more difficult, both because of its higher dimensionality, but also because it adds the state inference problem on top of control problem. In a lot of cases a problem which reinforcement learning algorithms are able to solve with direct state access is unsolvable with only pixel-based observations. In the cases where it is possible, the training time is much longer because much more samples are required. This is problematic because reinforcement learning is rather inefficient as it is. The high number of required samples in particular prohibits its use outside of simulated environments, while learning on agents in the real world is the ultimate goal in most practical applications.

1.4 Efforts to make reinforcement learning more efficient

1.4.1 Utilizing a world model

An important classification of reinforcement learning algorithm is the one between model-based and model-free algorithms. As the name suggests, model-free algorithms do not form an explicit model of the environment. Instead, they function as black-box optimization algorithms, simply finding actions which maximize reward without other concerns such as predicting the states resulting from those actions. In other words, they only predict the reward of actions in given states. Model-based algorithms on the other hand learn an explicit model of the environment and use it to plan their actions. Thus, they learn the dynamics of the environment and use that knowledge to choose actions which lead the agent to states with high reward. Both classes have their benefits and their drawbacks. Since model-free algorithms do not require any knowledge of environment dynamics to operate, they are more widely applicable and usually achieve better performance. But the fact that they can not leverage environment dynamics to create plans results in a harder learning problem: they need to implicitly learn those dynamics while only being provided the reward signal. This makes them much less sample-efficient.

Unfortunately, the model-based twin learning objective of learning the best action-choosing policy to maximize the reward over time, and the learning of the model, results in fundamental training instabilities which usually results in worse final performance. In simple terms, the reason behind this is the following one: in the beginning of the learning process, both the policy and the model perform poorly. For the model to perform better, the agent needs to explore the environment and update its model. However, many parts of the environment are inaccessible to a poorly performing agent: for example, if an agent is playing a computer game, and it is not able to progress to further sections of the game, it will not be able to construct a model of that portion of the game. Thus, to explore the environment and improve its model, it needs to first learn exploit the model and perform sufficiently well using it. Furthermore, what it learned at this stage may become obsolete as the model changes. How bad this problem is depends on the specifics of the setting,

and there are many ways to ameliorate it, but in most cases the necessary trade-offs result in a lower final performance.

1.4.2   Utilizing state representations

In the case of image-based observations, an alternative to using model-based methods is to additionally deal only with the problem of extracting states from images. This can be done by using auxiliary unsupervised representation learning goals. More concretely, the idea is to utilize learning signals other than the reward signal to make the model-free learning more sample-efficient. In particular, this amounts to learning a latent representation of the environment, i.e. finding a lower-dimensional embedding of the observations, and learning a policy in this space.
The benefits of this approach are two-fold:
- It is know that, in general, lower-dimensional optimization problems are easier to solve than higher-dimensional ones.
- Empirical findings show that when reinforcement algorithms have direct state access, they learn much faster and often achieve better final results.

Since inferring states from observations is not directly related to the reward, unsupervised learning techniques should aid in feature extraction and thus make learning more sample-efficient. This is the approach investigated in this thesis. Since the goal is not to learn the dynamics of the environment, but simply to find an equivalent, but lower-dimensional representation of it, it is this approach should not suffer from the problems faced by model-based approaches.

1.5   Goal of the thesis

Given the previous discussion, the goal of the thesis may be presented: the idea is to investigate how unsupervised learning techniques can be combined with model-free algorithms in order to increase their sample-efficiency. To make this a concrete and manageable goal, we constrain ourselves to the problem of learning from images and to problems with discrete action spaces in particular. To be able to compare our results to those of other researchers, we will test our algorithms on the standard benchmark tasks in the field, namely Atari-57 games [Bel+13a]. Of course, we are not the first to suggest such an approach. An overview of the field is provided in 4.

1.5.1   Hypothesis

As already stated, we believe that leveraging unsupervised learning techniques to learn state representations will make reinforcement learning from images more sample-efficient. Testing whether this is in fact true is one of our tasks. As will be shown in 3.2, there are many different unsupervised learning techniques which can be adapted to the goal of state representation learning. Furthermore, there are many different ways in which state representation learning can be integrated in the reinforcement learning process. In this thesis, our goal is not to arrive at a new state-of-the-art algorithm, but to investigate which properties of both the state representation learning and its integration with reinforcement learning yield better results. We will not

test all of the existing approaches, but rather identify their common properties, form hypothesis based on those properties and perform tests on a simple implementation. Here we offer our hypothesis:

1. State representations found by general unsupervised learning techniques will not equate to true states, although they will be closer to them than raw image-based observations. Reinforcement learning algorithms are able to implicitly learn true states, but because they do so indirectly and by using the weak reward signal, they do so very slowly. Thus we hypothesize that allowing the reinforcement learning algorithm to continue updating feature extraction provided the state representation learning algorithm will perform better than feature extraction learned solely through state representation learning. We further hypothesize that the best feature extraction will be obtained if both state representation learning algorithm and the reinforcement learning algorithm continuously update the feature extractor throughout the entire training process.

2. We hypothesize that state representation learning algorithms whose learned features better match the underlying Markov decision process will yield better results. This for example means that representations learned on future prediction tasks will perform better than those which are not incentivized to learn dynamics.

3. Finally, we hypothesize that strong regularization of the state representation learning algorithms will yield better results. We believe that proper regularization will yield broader features which the reinforcement learning algorithm will more easily integrate with.

1.5.2   Contributions

As already mentioned, our main goal is not to produce a state-of-the-art algorithm, but rather to find and investigate general properties which state representation learning algorithms should have and how they should best be integrated with reinforcement learning algorithms. In this thesis we provide the following contributions:

1. A systematic overview of recent works which leverage state representation learning to make model-free reinforcement learning more sample-efficient.

2. Extensive testing of our hypothesis which illuminate the problem and pave the wave for further algorithm development.

3. Implementation of our method in a high-quality reinforcement learning library. Despite the fact that our method is not the best available one, its generality and its implementation makes it easily accessible to practitioners and helps researchers who wish to build on top of it.

1.6   Outline

The rest of this text is organized as follows. We begin by describing the basics of reinforcement learning in 2. Here the problem setting and basic concepts are covered. The main classes of reinforcement learning are introduced in 2.3. Having the basics of reinforcement learning established, in 2.4 we introduce the reinforcement

learning algorithm we use in our implementation and discuss common reinforcement learning problems in 2.5. Following the reinforcement learning discussion, we turn our attention to unsupervised representation learning on images in 3.1 and discuss common state representation learning approaches for control problems in 3.2. Following the background, discuss related work in 4. Here we cover several existing approaches to bolstering the sample-efficiency of model-free reinforcement learning with state representation learning. Having covered the field, in 5.1 we identify key factors which lead to state representations which can be leveraged by reinforcement learning algorithms. In other words, we then form the basis for our hypothesis.

2

Reinforcement learning

2.1 Problem setting

In the usual engineering approach to problems, prior scientific knowledge is used to first describe the problem and then to define it mathematically. Once this is done, unknown variables are measured and solutions are calculated. This approach works if the inherent stochasticity of the environment can be controlled, i.e. if bounds of stochasticity are known the solution account for them and be designed to be robust to them. But some problems have circumstances which can not be known in advance, or which are incredibly hard to hand-engineer.

In those cases, an entirely different approach becomes the only viable one: designing a system which can produce and refine its own solution, or in other words, designing a system which, in a way, learn the solution by itself. This is the idea behind the learning-based approaches: automating the process of learning. Crucially, now the world and how it operates is unknown and has to be discovered. The schematic 2.1 shows how this process is formulated in reinforcement learning. Reinforcement

action

agent environment

observation, reward

Figure 2.1: Conceptual schematic of reinforcement learning.

learning is a 2-step iterative process. The **agent**, which represents the computer program, takes **actions** in its **environment**. It then **observes** the resulting state of the environment and is also given a **reward** which is a function mapping every state of the environment to a number.

To help introduce reinforcement learning formally, first the simplest possible problem to which reinforcement learning is the best solution is described.

2.1.1 Bandit problems

Consider the following problem. The agent is faced with k different gambling slot machines. Each of them give random rewards under an unknown distribution. At

each turn, the agents has to select one of the machines and pull its lever. The goal is to maximize the expected total reward over some number of turns. If the agent knew the distribution of rewards of each of the slot machines, it would simply choose the one with the highest expected reward. However, the agent does not have access to that information and hence it can not effectively **exploit** the environment to obtain the highest rewards. Instead, it is forced to **explore** the environment in order to learn the probability distribution of the reward. In this problem, this amounts to pulling different levers and recording the received payoffs.

The key issue is thus to balance exploitation and exploration. It is compounded by the fact that the agent is given only a finite amount of time, or a finite number of lever pulls in this case. While one could leverage Bayesian statistics to construct an optimal solution to this simplest formulation of the bandit problem, this quickly becomes intractable as the complexity of the problem is increased. Namely, the described problem is stationary, in the sense that past actions do not influence the state of the world: the slot machines do not change as different levers are pulled. Said another way, the k-bandit problem has a single **state**. This is of course not the case in most problems of interest. To model the agent's effect on the environment, additional structure needs to be introduced. This is done in the following sections.

2.1.2 Markov Decision Processes

To model environments in which states change, Markov chains are used. They capture the stochastic nature of state transitions, while Markov property allows for easier mathematical analysis. The schematic of a Markov chain is shown in 2.2.

Figure 2.2: Schematic of a Markov chain.

Formally, a Markov chain $\mathcal{M}$ is defined by its state space $\mathcal{S}$ with discrete or continuous state $\boldsymbol{s} \in \mathcal{S}$ and the transition operator $\mathcal{T}$. The notation $\boldsymbol{s}_t$ denotes the state at time t and it is a vector of real numbers. The transition operator allows for a succinct description of environment dynamics. For a transition probability $p(s_{t+1}|s_t)$, let $\mu_{t,i} = p(\boldsymbol{s}_t = i)$ and $\mathcal{T}_{i,j} = p(\boldsymbol{s}_{t+1} = i|\boldsymbol{s}_t = j)$. Then $\overrightarrow{\mu}_t$ is a vector of probabilities and $\overrightarrow{\mu}_{t+1} = \mathcal{T}\overrightarrow{\mu}_t$. Importantly, $\mathcal{T}$ is linear.

To model the agent's actions, actions are added as priors to state transition probabilities in the Markov chain. With this and the definition of the reward function, a Markov decision process is constructed. It's schematic can be seen in 2.3.

Figure 2.3: Schematic of a Markov decision process.

The Markov decision process is thus defined by the tuple $\mathcal{M} = \{\mathcal{S}, \mathcal{A}, \mathcal{T}, r\}$. $\mathcal{A}$ denotes the action space, where $\boldsymbol{a} \in \mathcal{A}$ is a continuous or discrete action and r is the reward function $r : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$. It should also be noted that the transition operator is now a tensor. Let $\mu_{t,j} = p(s_t = j), \xi_{t,k} = p(a_t = k), \mathcal{T}_{i,j,k} = p(s_{t+1} = i | s_t = j, a_t = k)$. Then $\mu_{t+1,i} = \sum_{j,k} \mathcal{T}_{i,j,k} \mu_{t,j} \xi_{t,k}$. Therefore, $\mathcal{T}$ retains its linearity.

Finally, partial observability also needs to be accounted for. To do so, a partially observable Markov decision process (POMDP) needs to be constructed. This is done by augmenting the Markov decision process to also include the observation space $\mathcal{O}$, where observations $\boldsymbol{o} \in \mathcal{O}$ denote the discrete or continuous observations and the emission probability $\mathcal{E}$ which describes the probability $p(\boldsymbol{o}_t | \boldsymbol{s}_t)$ of getting the observation $\boldsymbol{o}_t$ when in state $\boldsymbol{s}_t$. The schematic can be seen in 2.4.

Figure 2.4: Schematic of a partially observable Markov decision process.

It is important to note that not all elements of POMDP are present in every problem: for example, the reward may be a deterministic function of the state and so on. In general through the text, to aid in simplifying notation, only the necessary elements will be explicity referenced in sketches and written out in the equations (most often using just the Markov decision process).

2.2 Key concepts in reinforcement learning

2.2.1 Policy

With the problem space being formally defined, definitions which will allow the construction of a reinforcement learning algorithm may be introduced. The reinforcement learning problem can be defined in finite or infinite time horizons. Different environments usually naturally fall in either category. In order for the agent to learn, it needs to be able to try out different actions from the same, or at least similar states. This is usually achieved by having the agent return to a set of starting states. The period between two such returns is called **an episode**. The agent selects actions based on its **policy** π. The policy is a function which maps states to actions. The schematic showing it in the context of a Markov decision process is given in 2.5.

The policy is a stochastic function. The intensity of stochasticity determines the trade-off between exploration and exploitation. To emphasize that the policy depends on some parameters θ, the notation π_θ is used.

Figure 2.5: Schematic of a Markov decision process with a policy π.

2.2.2 Goal of reinforcement learning

For simpler notation, the finite horizon form is assumed for the following definitions. Since the environment is modeled as a Markov decision process, the probability of observing a trajectory of states and actions can be written as:

$$\underbrace{p_\theta(\boldsymbol{s}_1, \boldsymbol{a}_1, \ldots, \boldsymbol{s}_T, \boldsymbol{a}_T)}_{p_\theta(\tau)} = p(\boldsymbol{s}_1) \prod_{t=1}^{T} \underbrace{\pi_\theta(\boldsymbol{a}_t|\boldsymbol{s}_t)p(\boldsymbol{s}_{t+1}|\boldsymbol{s}_t, \boldsymbol{a}_t)}_{\text{Markov chain on } (\boldsymbol{s},\boldsymbol{a})} \tag{2.1}$$

A bit more explicitly, we can a transition probability as:

$$p((\boldsymbol{s}_{t+1}, \boldsymbol{a}_{t+1})|(\boldsymbol{s}_t, \boldsymbol{a}_t)) = p((\boldsymbol{s}_{t+1}|(\boldsymbol{s}_t, \boldsymbol{a}_t))\pi_\theta(\boldsymbol{a}_{t+1}|\boldsymbol{s}_{t+1}) \tag{2.2}$$

With this, a formal definition of the goal of reinforcement learning can be given. It is to find policy parameters $\theta^\star$ such that:

$$\theta^\star = \operatorname*{argmax}_\theta \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[\sum_t r(\boldsymbol{s}_t, \boldsymbol{a}_t) \right] \tag{2.3}$$

$$= \operatorname*{argmax}_\theta \sum_t^T \mathbb{E}_{(\boldsymbol{s}_t, \boldsymbol{a}_t) \sim p_\theta(\boldsymbol{s}_t, \boldsymbol{a}_t)} \left[r(\boldsymbol{s}_t, \boldsymbol{a}_t) \right] \tag{2.4}$$

To ensure that the expected sum of rewards, also know as the **return**, is finite in the infinite horizon case, a **discount factor** $0 < \gamma < 1$ is introduced in the sum. The discount factor also plays a role in modelling because usually it makes sense to value immediate rewards more. It is important to note that the *expected* sum of rewards is maximized . This makes the goal a smooth and differentiable function of the parameters, which means gradient descent can be employed to find the optimal parameters. This leads us to the first class of reinforcement learning algorithms: policy gradient algorithms. They will be introduced with the other classes of algorithm in 2.3.1, while additional concepts required by other classes of algorithms will be introduced in the following subsection.

2.2.3 Value functions

Value functions are functions which map states or state-action pairs to the expected returns obtained under a fixed policy. They are a concept from dynamic programming. In fact, reinforcement learning can be interpreted as an extension of dynamic programming, as shall be done in the following subsection. Having that said, value

function can be interpreted in other ways as well. The **Q-function** maps state-action pairs to the estimated sum of returns under policy π_θ:

$$Q^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t) = \sum_{t'=t}^{T} \mathbb{E}_{\pi_\theta} \left[r(\boldsymbol{s}_{t'}, \boldsymbol{a}_{t'}) | \boldsymbol{s}_t, \boldsymbol{a}_t \right] \tag{2.5}$$

thus denoting the expectred total reward from taking $\boldsymbol{a}_t$ in $\boldsymbol{s}_t$. **Value functions** map states to to the expected sum of rewards under policy π_θ:

$$V^\pi(\boldsymbol{s}_t) = \sum_{t'=t}^{T} \mathbb{E}_{\pi_\theta} \left[r(\boldsymbol{s}_{t'}, \boldsymbol{a}_{t'} | \boldsymbol{s}_t) \right] \tag{2.6}$$

The connection between the two is the following:

$$V^\pi(\boldsymbol{s}_t) = \mathbb{E}_{\boldsymbol{a}_t \sim \pi(\boldsymbol{s}_t, \boldsymbol{a}_t)} \left[Q^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t) \right] \tag{2.7}$$

With these definitions, the reinforcement learning objective may also be written as:

$$\mathbb{E}_{\boldsymbol{s}_1 \sim p(\boldsymbol{s}_1)} \left[V^\pi(\boldsymbol{s}_1) \right] \tag{2.8}$$

2.3 Classes of reinforcement learning algorithms

2.3.1 Policy gradients

Policy gradients are derived by directly solving for the reinforcement learning objective using gradient descent, where the derivative is taken with respect to the policy parameters. To do so, the reinforcement learning objective needs to be evaluated. To make this a bit easier to follow, we introduce a notational shorthand:

$$\theta^\star = \underset{\theta}{\arg\max} \underbrace{\mathbb{E}_{\tau \sim p_\theta(\tau)} \left[\sum_t r(\boldsymbol{s}_t, \boldsymbol{a}_t) \right]}_{J(\theta)} \tag{2.9}$$

$J(\theta)$ is estimated by making rollouts from the policy. Simply put, the agent collects experience under the current policy, and the average return is used as the estimate. In the equation below i is the sample index. i, t denotes the t^{th} timestep in the i^{th} sample:

$$J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[\sum_t r(\boldsymbol{s}_t, \boldsymbol{a}_t) \right] \approx \frac{1}{N} \sum_i \sum_t r(\boldsymbol{s}_{i,t}, \boldsymbol{a}_{i,t}) \tag{2.10}$$

Simplifying the notation further, we get:

$$J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} \underbrace{[r(\tau)]}_{\sum_{t=1}^{T} r(\boldsymbol{s}_t, \boldsymbol{a}_t)} = \int p_\theta(\tau) r(\tau) \, d\tau \tag{2.11}$$

The goal now is to compute the derivative of the estimated reinforcement learning objective:

$$\nabla_\theta J(\theta) = \int \nabla_\theta p_\theta(\tau) r(\tau) \, d\tau \tag{2.12}$$

Since the goal of this text is just to introduce the necessary concepts and algorithms, the derivation(s) will be ommited. We encourage the interested reader to consult the literature [SB18; Lev21] to find more information.

Here we will just note that it is crucial that the final expression can be estimated by sampling the agent's experience as the other quantities are not available. The resulting expression for the policy gradient is:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[\left(\sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(\boldsymbol{a}_t | \boldsymbol{s}_t) \right) \left(\sum_{t=1}^{T} r(\boldsymbol{s}_t, \boldsymbol{a}_t) \right) \right] \tag{2.13}$$

It can be evaluated through sampling:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} \left(\sum_{t=1}^{T} \nabla_\theta \log \pi_\theta(\boldsymbol{a}_{i,t} | \boldsymbol{s}_{i,t}) \right) \left(\sum_{t=1}^{T} r(\boldsymbol{s}_{i,t}, \boldsymbol{a}_{i,t}) \right) \tag{2.14}$$

The estimated gradient can be used to perform gradient ascent. This is the backbone of the REINFORCE algorithm, also known as "vanilla policy gradient":

REINFORCE algorithm:
1. sample $\{\tau^i\}$ from $\pi_\theta(\boldsymbol{a}_t | \boldsymbol{s}_t)$ by running the policy
2. use the samples to estimate the gradient of the objective:
$\nabla_\theta J(\theta) \approx \sum_i \left(\sum_t^T \nabla_\theta \log \pi_\theta(\boldsymbol{a}_{i,t} | \boldsymbol{s}_{i,t}) \right) \left(\sum_t r(\boldsymbol{s}_{i,t}, \boldsymbol{a}_{i,t}) \right)$
3. update the policy function by performing a step of gradient ascent:
$\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

This algorithm does not work well in practice. The main reason for that is that the variance of returns is very high. However, there are a number of modifications which dramatically improve its performance. Since the goal of this text is not to outline every reinforcement learning algorithm, we will introduce only the modifications which outline general trade-offs and principles in reinforcement learning algorithm design.

2.3.1.1   Baselines

The policy gradient in the REINFORCE algorithm lacks some important properties. One of them is that it should, ideally, make bad actions less likely and good actions more likely. However, if all rewards are positive, then all actions' probabilities will be increased, only by different amounts. This can be changed if a **baseline** b is added to actions:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} \nabla_\theta \log p_\theta(\tau)[r(\tau) - b] \tag{2.15}$$

$$b = \frac{1}{N} \sum_{i=1}^{N} r(\tau) \tag{2.16}$$

This addition does not change the gradient in expectation, i.e. it does not introduce bias, but it does change its variance. Although an optimal bias can be calculated, it is rarely used in practice due to its computational cost. Using baselines is one of the key ideas in actor-critic algorithms so they will be discussed further in 2.3.2.

2.3.1.2 Off-policy gradients

An important property of the REINFORCE algorithm is that it is an **on-policy** algorithm. This means that new samples need to be collected for every gradient step. The reason behind this is the fact that the expectation of the gradient of the return needs to be calculated with respect to the current parameters of the policy. In other words, because the policy changes with each gradient step, old samples are effectively collected under a different policy. This means that they can not be used to calculate the expected gradient of the return with respect to the current policy as it would not produce those rollouts:

$$\nabla_\theta J(\theta) = \underbrace{\mathbb{E}_{\tau \sim p_\theta(\tau)}}_{\text{this is the trouble!}} [\nabla_\theta p_\theta(\tau) r(\tau)] \tag{2.17}$$

If the policy is a neural network, which requires small gradient steps, the cost of generating a large number of samples for every update could make the algorithm entirely infeasible. This of course depends on the cost of generating samples, which is entirely problem dependent — policy gradient algorithms are often the best solution when the cost of generating samples is low.

However, on-policy algorithms can be turned into off-policy algorithms through **importance sampling**, which is the name given to the following mathematical identity:

$$E_{x \sim p(x)}[f(x)] = \int p(x) f(x)\, dx \tag{2.18}$$

$$= \int \frac{q(x)}{q(x)} p(x) f(x)\, dx \tag{2.19}$$

$$= \int q(x) \frac{p(x)}{q(x)} f(x)\, dx \tag{2.20}$$

$$= E_{x \sim p(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \tag{2.21}$$

which is exact in expectation. To use importance sampling to create an off-policy policy gradient algorithm, certain approximations need to be made. Again, the details of the derivation are omitted and what follows is just the final result.

$$\nabla_{\theta'} J(\theta') \approx \frac{1}{N} \sum_{i=1}^{N} \sum_{t=1}^{T} \frac{\pi_{\theta'}(\boldsymbol{a}_{i,t}|\boldsymbol{s}_{i,t})}{\pi_\theta(\boldsymbol{a}_{i,t}|\boldsymbol{s}_{i,t})} \nabla_{\theta'} \log \pi_{\theta'}(\boldsymbol{s}_{i,t}, \boldsymbol{a}_{i,t}) \hat{Q}_{i,t} \tag{2.22}$$

To get this equation, the factor $\frac{\pi_{\theta'}(\boldsymbol{s}_{i,t})}{\pi_\theta(\boldsymbol{s}_{i,t})}$ had to be ignored in the expression because it is impossible to calculate the state marginal probabilities. This means that the expression works only if $\pi_{\theta'}$ is not too different from π_θ. The justification for this can be found in the previously referenced literature. What is important in the context of the thesis is that we will be interested only in off-policy methods as they are inherently more sample-efficient and out goal is to increase the sample-efficiency of reinforcement learning.

2.3.1.3   Advanced policy gradients

The basic algorithm we have outlined is essentially just a basic gradient descent method. From convex optimization, we know that it can be made much better if second order derivatives or their approximations are used. For example, conjugate gradient descent can be used. Further, there are various ways in which this optimization problem can be better conditioned. Such improvements led to algorithms such as PPO and TRPO, which will not be discussed here.

2.3.2   Actor-critic algorithms

Actor-critic methods can be seen as making a different trade-off between variance and bias in policy gradient estimation. Consider with the following observation: [1]

$$\nabla_\theta J(\theta) \approx \frac{1}{N}\sum_{i=1}^{N}\sum_{t=1}^{T}\nabla_\theta \log \pi_\theta(\boldsymbol{a}_{i,t}|\boldsymbol{s}_{i,t})\underbrace{\left(\sum_{t'=t}^{T}r(\boldsymbol{s}_{i,t},\boldsymbol{a}_{i,t})\right)}_{\hat{Q}_{i,t}:\ \text{``reward to go''}} \tag{2.23}$$

Simply put, in the policy gradient method a single-run Monte-Carlo (MC) estimate is used to estimate the return. This causes high variance, while incurring no bias. Another option is to try to estimate the full expectation $\hat{Q}_{i,t} \approx \sum_{t'=t}^{T}\mathbb{E}_{\pi_\theta}\left[r(\boldsymbol{s}_{t'},\boldsymbol{a}_{t'})|\boldsymbol{s}_t,\boldsymbol{a}_t\right]$. Since the estimate won't be perfect, it will introduce bias. Of course, using multiple runs from the same state-action pair would reduce variance, but this is sometimes impossible to procure and is certainly more costly. However, if the "reward to go" estimator can generalize between states, it will be able to produce good estimates regardless.

Like the policy, the return estimator will have to be learned. In this approach, the policy is also called the **actor** and the return estimator is called the **critic**. If the correct Q-function was available (i.e. not the estimate, but the actual values), the policy gradient estimate could be improved by using it both to estimate the return and as a baseline:

$$\nabla_\theta J(\theta) \approx \frac{1}{N}\sum_{i=1}^{N}\sum_{t=1}^{T}\nabla_\theta \log \pi_\theta(\boldsymbol{a}_{i,t}|\boldsymbol{s}_{i,t})(Q(\boldsymbol{s}_{i,t},\boldsymbol{a}_{i,t})-b) \tag{2.24}$$

$$b_t = \frac{1}{N}\sum_i Q(\boldsymbol{s}_{i,t},\boldsymbol{a}_{i,t}) \tag{2.25}$$

However, having a baseline that depends on actions leads to bias. Thus a state-dependent baseline is employed:

$$V(\boldsymbol{s}_t) = E_{\boldsymbol{a}_t \sim \pi_\theta(\boldsymbol{s}_t,\boldsymbol{a}_t)}[Q(\boldsymbol{s}_t,\boldsymbol{a}_t)] \tag{2.26}$$

Since the value function 2.6 gives the expected return of the average action, it is possible to calculate how much better a certain action is by substracting its Q-value 2.5 for the value function. The result is called the **advantage function**:

$$A^\pi(\boldsymbol{s}_t,\boldsymbol{a}_t) = Q^\pi(\boldsymbol{s}_t,\boldsymbol{a}_t) - V^\pi(\boldsymbol{s}_t) \tag{2.27}$$

---

[1] In this equation, the summation of rewards is done from time t to T because actions and states prior to that time do not affect the return from that time onward. This leveraging of causality reduces the variance of the estimate.

Thus, either the Q-function, the value function or the advantage function can be learned. Of these, it is best to learn the value function because there are less states than state-action pairs. The advantage function is then calculated in the following way:

$$A^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t) \approx r(\boldsymbol{s}_t, \boldsymbol{a}_t) + V^\pi(\boldsymbol{s}_{t+1}) - V^\pi(\boldsymbol{s}_t) \tag{2.28}$$

The value function can be estimated through samples:

$$V^\pi(\boldsymbol{s}_t) \approx \sum_{t'=t}^{T} r(\boldsymbol{s}_{t'}, \boldsymbol{a}_{t'}) \tag{2.29}$$

After collecting many such samples

$$\left\{ \left(\boldsymbol{s}_{i,t}, \underbrace{\sum_{t'=t}^{T} r(\boldsymbol{s}_{i,t}, \boldsymbol{a}_{i,t})}_{y_{i,t}} \right) \right\} \tag{2.30}$$

the value function can be fitted through supervised regression with the following loss:

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_i ||\hat{V}_\phi^\pi(\boldsymbol{s}_i) - y_i||^2 \tag{2.31}$$

This process can be sped up with bootstrapped estimates:

$$y_{i,t} = \sum_{t'=t}^{T} E_{\pi_\theta} \left[r(\boldsymbol{s}_{t'}, \boldsymbol{a}_{t'}) | \boldsymbol{s}_{i,t} \right] + V^\pi(\boldsymbol{s}_{i,t+1}) \approx r(\boldsymbol{s}_{i,t}, \boldsymbol{a}_{i,t}) + \hat{V}_\phi^\pi(\boldsymbol{s}_{i,t+1}) \tag{2.32}$$

This will further reduce variance, but again increase bias.
Fortunately, the trade-off between bias and variance can be tuned. In the Monte Carlo estimate, the entire trajectory was used to estimate the return. In the bootstrap estimate, only a single step in the future was used along with the estimate. Instead, a **n-step** return estimator can used:

$$\hat{A}_n^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(\boldsymbol{s}_{t'}, \boldsymbol{a}_{t'}) - \hat{V}_\theta^\pi(\boldsymbol{s}_t) + \gamma^n \hat{V}_\theta^\pi(\boldsymbol{s}_{t+n}) \tag{2.33}$$

In most cases the ideal trade-off for n lies somewhere between 1 and ∞ (the MC estimate). Finally, an average of all n-step return estimators can be used. This is called the generalized advantage estimator (GAE):

$$\hat{A}_{GAE}^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t) = \sum_{n=1}^{\infty} (\gamma\lambda)^{t'-t} r(\boldsymbol{s}_{t'}, \boldsymbol{a}_{t'}) + \gamma \hat{V}_\theta^\pi(\boldsymbol{s}_{t'+1}) - \hat{V}_\theta^\pi(\boldsymbol{s}_{t'}) \tag{2.34}$$

where the factor λ controls the weight of future values.
Combining this into an iterative algorithm, and fixing the issues of naiive implementations results in the following algorithm:

Actor-critic algorithm template
1. take action $\boldsymbol{a} \sim \pi_\theta(\boldsymbol{a}|\boldsymbol{s})$, observe transition $(\boldsymbol{s}, \boldsymbol{a}, \boldsymbol{s}', r)$ and store it in the replay buffer $\mathcal{R}$
2. sample a batch $\{(\boldsymbol{s}_i, \boldsymbol{a}_i, \boldsymbol{s}'_i, r_i)\}$ from buffer $\mathcal{R}$
3. update the Q-value estimator $\hat{Q}_\theta^\pi$ by using the target:
   $y_i = r_i + \gamma \hat{Q}_\theta^\pi(\boldsymbol{s}'_i, \boldsymbol{a}'_i) \forall \boldsymbol{s}_i, \boldsymbol{a}_i$
4. compute the policy gradient estimate with:
   $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\boldsymbol{a}_i^\pi|\boldsymbol{s}_i)\hat{Q}^\pi(\boldsymbol{s}_i, \boldsymbol{a}_i^\pi)$, where $\boldsymbol{a}_i^\pi \sim \pi_\theta(\boldsymbol{a}|\boldsymbol{s}_i)$
5. update the policy function by performing a gradient step:
   $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

2.3.3   Value function methods

Value function methods use only the critic from actor-critic algorithms. Suppose that the advantage function $A^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t)$ is known. It tells us how much better the action $\boldsymbol{a}_t$ is than the average action according to the policy π. Thus, if provided with the advantage function, a deterministic **greedy policy** could be construct:

$$\pi_{\text{greedy}}(\boldsymbol{s}_t|\boldsymbol{a}_t) = \begin{cases} 1 & \text{if } \boldsymbol{a}_t = \text{argmax}_{\boldsymbol{a}_t} A^\pi(\boldsymbol{s}_t, \boldsymbol{a}_t) \\ 0 & \text{otherwise} \end{cases} \tag{2.35}$$

which would yield the highest expected return. In other words, if the advantage function is known, the policy would be reduced to the argmax operation. This approach has roots in dynamic programming.

2.3.3.1   Dynamic programming

Dynamic programming refers to a collection of algorithms that can be used to compute optimal policies given a perfect model of the environment. They are of limited utility in reinforcement learning due to the perfect model requirement and their great computational expense, but are important theoretically: they provide an essential foundation for understanding the other methods. Usually a finite Markov decision process is assumed. Dynamic programming can be applied to continuous problems as well, but exact solution exist only in special cases.
Value iteration and Q-learning make up two fundamental algorithms of Reinforcement Learning.Q learning is an Off policy algorithm, which means it uses a different policy for exploring actions from the target policy of being optimal. Many of the amazing feats in RL over the past decade, such as Deep Q-Learning for Atari, or AlphaGo, Rainbow were rooted in these foundations. As stated in section 2.2.3 value function ,is a measure of the expected reward you can receive from any given state given an MDP and a policy describing which actions an agent takes in each state. Value iteration is a computational algorithm that provides a means of finding the optimal policy π^* . The algorithm works by iteratively determining the value of being in each state, assuming that the agent takes the best possible action in that state under the current estimate of the value function.Fitted value iteration algorithm are used for approximating the value function of a continuous state MDP.Unlike value iteration over discrete state of states fitted value iteration can not

always converge.In order for certain dynamic programming algorithms (e.g. policy iteration, value iteration) converge to a unique fixed point Bellman equation was rewritten as operator.

2.4 Deep Q-networks and their extensions

The success of a deep neural network in computer vision has introduced a new paradigm in learning from raw pixels. The goal of Deep Reinforcement Learning is to connect the reinforcement learning algorithm to a deep neural network that operates directly on RGB images and efficiently processes training data using stochastic gradient updates [Mni+13].

The use of label data and the assumption that the distribution of data is identical and independent throughout the deep learning training process makes it complex to use it directly in reinforcement learning algorithms, which must be able to learn from sparse,noisy,delayed reward and highly correlated data. To address these issues and successfully apply deep learning to reinforcement learning [Mni+13] used the experience replay mechanism [Mni+15] which randomly samples previous transitions and thereby smooths the distribution of training over many past behaviors. A convolutional neural network was used to learn successful control policies from raw video data in a complex reinforcement learning environment. The network is trained with a variant of Q-learning algorithm,with stochastic gradient descent to update weights.

This architecture was based on the Tesauro TD-Gammon architecture [17] that updates the parameters of the network that estimate the value function directly from the sampled experience of the policy. Similarly to this approach, the online network in DQN uses a technique known as experience replay [18] where the agent's experiences at each time-step, $]_t = (s_t, a_t, r_t, s_t + 1)$ is stored in a data set $\mathcal{D} = e1, ..., e_N$ pooled over many episodes in a replay memory. By drawing random samples from this pool of stored experiences, the Q-learning is updated.After performing the experience replay, the agent selects and executes an action according to a ϵ-greedy policy. The use of experience replay and target networks enables relatively stable learning of Q values, and led to super-human performance on several Atari games. The advantage of using deep Q-learning over Q-learning includes allowing to have greater sample efficiency, reduced variance by randomizing the sample bias, and avoiding being stuck in local minimum. The drawbacks of deep Q-learning is that it only handle discrete, low-dimensional action spaces. Although the initial architecture of deep Q learning introduced by [Mni+13] paves the way for the use of deep neural network in reinforcement learning, it comes with its own drawbacks;this leads to more studies and an improved DQN architecture.In this section, we will discuss these extensions of DQN architectures.

2.4.1 Double Deep Q-networks: DDQN

DQN suffer from overestimation bias due to due to the maximization step in optimisation function in Q-learning.Q-learning can overestimate actions that have been

tried often and the estimations can be higher than any realistic optimistic esti-
mate .Double Q-learning [19], addresses this overestimation by decoupling, in the
maximization performed for the bootstrap target, the selection of the action from
its evaluation.Double Q-learning stores two Q-functions,The average of the two Q
values for each action and then performed ϵ-greedy exploration with the resulting
average Q values.It was successfully combined with DQN to reduce overestimations.

2.4.2 Prioritized replay

The main use of replay buffer is to sample transitions with maximum probabil-
ity.Both DQN and DDQN samples experiences uniformly.Prioritized replay [20]
samples transitions using the maximum priority ,providing a bias towards recent
transitions and stochastic transitions even when there is little left to learn about
them.

2.4.3 Dueling Network

Dueling Network was designed for value based learning,this architecture separates
the representation of sate-value and state-dependent action advantages without su-
pervision[6].its consists of two streams that represents the value and advantage func-
tions,while sharing a common convolutional feature learning module.This network
has a single Q-learning network with two streams that replace DQN architecture[3].

$$Q(s, a; \theta, \alpha, \beta) = V(s, \theta, \beta) + A(s, a; \theta, \alpha) \tag{2.36}$$

2.4.4 Multi-step learning

Previously stated extension of DQN have indicated that the use deep learning has
enhanced the learning capability of Q-learning.The performance of Q-learning is still
limited by greedy action selection after accumulating a single reward.An alternative
approach was multi-step targets:

$$y_{j,t} = \sum_{t'=t}^{t+N-1} \gamma^{t-t'} r_{j,t'} + \gamma^N \max_{\boldsymbol{a}_{j,t+N}} Q_{\phi'}(\boldsymbol{s}_{j,t+N}, \boldsymbol{a}_{j,t+N}) \tag{2.37}$$

A multi-step variant of DQN is then defined by minimizing the alternative loss[16],

$$R_t^{(n)} + \gamma_t^{(n)} \max_{a'} q_{\theta}^-(S_{t+n}, a') - q_{\theta}(S_t, A_t) \tag{2.38}$$

2.4.5 Noisy Nets

The one limitation of ϵ-greedy policy is many actions must be executed to collect the
first reward.Noisy Nets proposed a noisy linear layer that combines a deterministic
and noisy stream.Depending on the learning rate the network ignores to learn the
noisy stream.

2.4.6 Integrated Agent:Rainbow

In the Rainbow architecture [Hes+18] several architecture changes included the one stated above where applied to DQN. Distributional loss was replaced by a multi-step variant. The target distribution was constructed by contracting the value distribution in St+n according to the cumulative discount, and shifting it by the truncated n-step discounted return. multi-step distributional loss with double Q-learning by using the greedy action in St+n selected according to the online network as the bootstrap action $a \cdot t + n$, and evaluating such action using the target network.

2.4.7 Deep autoencoders

Reinforcement learning requires learning from large high-dimensional image data-set.For example, In Atari games the environment is composed of images with 210 * 160 pixels and 128 color palette. Each image is made up of hundreds of pixels, so each data point has hundreds of dimensions. The manifold hypothesis states that real-world high-dimensional data actually consists of low-dimensional data that is embedded in the high-dimensional space. This is the motivation behind dimensionality reduction techniques, which try to take high-dimensional data and project it onto a lower-dimensional surface.

Autoencoders are a special kind of neural network used to perform dimensionality reduction. They act as an identity function, such that an auto encoder learns to output whatever is the input.They are composed of two networks, an encoder e and a decoder d.

The encoder learns a non-linear transformation that projects the data from the original high-dimensional input space X to a lower-dimensional latent space Z. This is called latent vector $z = e(x)$. A latent vector is a low-dimensional representation of a data point that contains information about x. This is commonly known as latent space representation, it contains all the important information needed to represent raw data points.Auto encoders manipulates the "closeness" of data in the latent space.

A decoder learns a non-linear transformation d:Z→X that projects the latent vectors back into the original high-dimensional input space X. This transformation takes the latent vector and reconstruct the original input data :

$$z = e(x) \rightarrow \hat{x} = d(z) = d(e(x)) \tag{2.39}$$

The autoencoder is trained to minimize the difference between the input x and the reconstruction $\hat{x}$ using a kind of reconstruction loss.

In traditional autoencoders, the latent vector should be easily decoded back to the original image as a result the latent space z can become disjoint and non-continuous. Variational autoencoders try to solve this problem.

In variational autoencoders, inputs are mapped to a probability distribution over latent vectors, and a latent vector is then sampled from that distribution. As a result the decoder becomes more robust at decoding latent vectors.

Specifically, instead of mapping the input x to a latent vector $z = e(x)$, we instead map it to a mean vector $\mu(x)$ and a vector of standard deviations $\sigma(x)$. These

parametrize a diagonal Gaussian distribution $\mathcal{N}(\mu, \sigma)$, from which we then sample a latent vector $z \sim \mathcal{N}(\mu, \sigma)$.

This is generally accomplished by replacing the last layer of a traditional autoencoder with two layers, each of which output $\mu(x)$ and $\sigma(x)$. An exponential activation is often added to $\sigma(x)$ to ensure the result is positive.

However, this does not completely solve the problem. There may still be gaps in the latent space because the outputted means may be significantly different and the standard deviations may be small. To reduce that, an auxiliary loss is added that penalizes the distribution $p(z|x)$ for being too far from the standard normal distribution $\mathcal{N}(\iota, \infty)$. This penalty term is the Kullback-Leibler(KL) divergence between $p(z|x)$ and $\mathcal{N}(\iota, \infty)$, which is given by$\mathbb{KL}\left(\mathcal{N}(\mu, \sigma) \parallel \mathcal{N}(0, 1)\right) = \sum_{x \in X}\left(\sigma^2 + \mu^2 - \log \sigma - \frac{1}{2}\right)$.This expression applies to two univariate Gaussian distributions by summing KL divergence for each dimension we are able to extend it to our diagonal Gaussian distributions.

This loss is useful for two reasons. First, we cannot train the encoder network by gradient descent without it, since gradients cannot flow through sampling (which is a non-differentiable operation). Second, by penalizing the KL divergence in this manner, we can encourage the latent vectors to occupy a more centralized and uniform location. In essence, we force the encoder to find latent vectors that approximately follow a standard Gaussian distribution that the decoder can then effectively decode.

2.5 Problems with RL

In previous studies presented in Chapter 3 we see some successful application of unsupervised learning techniques applied to improve the performance of the underlying RL algorithms, even though these and other studies conducted previously have shown remarkable results. RL comes with the following challenges.

One of the most difficult aspects of RL is learning efficiently with little data. The term "sample efficiency" refers to an algorithm that makes the most of a given sample. To put it another way, it's the amount of experience the algorithm has to gain during training in order to achieve efficient performance. The difficulty is that the RL system takes a long time to become efficient.

Neural networks are opaque black boxes whose workings are mysteries to even the creators. They are also increasing in size and complexity, backed by huge data sets, computing power and hours of training.This referred to as Reproducibility crisis, These factors make RL models very difficult to replicate.

Another major challenge in Rl is agents are trained in simulated environment ;In this environment they can fail and learn, but they do not have the opportunity to fail and learn in real-life scenarios. Usually, in real environments, the agent lacks the space to observe the environment well enough to use past training data to decide on a winning strategy. This also includes the reality gap, where the agent cannot gauge the difference between the learning simulation and the real world.

The reward technique discussed in the previous sections is not foolproof. Since the

rewards are sparsely distributed in the environment, a possible issue is an agent not observing the situation enough to notice the reward signals and maximise specific actions. This also occurs when the environment cannot provide reward signals in time; for instance, in many situations, the agent receives a green flag only when it is close enough to the target.

Curiosity-driven methods are widely used to encourage the agent to explore the environment and learn to tackle tasks in it. The researchers in the paper 'Curiosity-driven exploration by self-supervised prediction' proposed an Intrinsic Curiosity Module (ICM) to support the agent in exploration and prompt it to choose actions based on reduced errors. Another approach is curriculum learning, where the agent is presented with various tasks in ascending order of complexity. This imitates the learning order of humans.

3

State representation learning

As discussed in 1.3, learning control from images is very desirable. Images, and observations in general, only implicitly provide information about the underlying state. Finding a good policy from observations, especially images, is much more difficult than finding a policy with direct state access because the state first needs to be inferred from those observations. Reinforcement learning algorithms can by themselves implicitly extract the relevant information from observations, but this at best results in much less sample-efficient training and at worst results in complete failure. Often a problem which a reinforcement learning algorithm can solve with direct state access, can not achieve any progress when provided only image observations.

Extracting lower-dimensional information from data in order to extract meaningful information from it is an established problem in machine learning broadly, and is referred to as **representation learning** in that context. Clearly, learning to extract stateful information from image observations can be viewed as a subset of representation learning and in this context it is referred to as **state representation learning**. One important aspect of representation learning is that the representations can be abstract and only implicitly represent the data in question. Such representations can be learned in an unsupervised manner: the goal is to learn latent representations which have useful features, and as such they can be directly optimized with regards to these features and not to conform to some explicit semantic form. The purpose of this chapter is to briefly discuss representation learning in general terms, and then to investigate how it can be applied to the problem of learning state representations from images for control problems. Importantly, the problem of learning a model which can be used to achieve control through planning will not be discussed, although there are similarities between the two.

The fundamental reason why this is a promising proposition is the fact that the learning signal generated from for example image reconstruction loss is substantially stronger than the reward signal, especially in settings with sparse rewards where it is not present most of the time. This thus means that the state representation will be learned quickly compared to policy learning. Hence they can be leveraged to aid the sample-efficiency in of reinforcement learning, despite their training starting at the same time as that of the policy.

3.1 Representation learning in general

In general, representation learning refers to the process of learning a parametric mapping from raw input data domain to a feature vector or tensor, in the hope of capturing and extracting more abstract and useful concepts that can improve performance of downstream tasks. This mapping should also meaningfully generalize well on new data. The following list, introduced in [**bengio2013representation**], summarizes different desirable characteristics of learned representations:

1. *smoothness*: f s.t. $x \approx y$ implies $f(x) \approx f(y)$

2. *multiple explanatory factors* a.k.a. disentangling features

3. *semi-supervised learning*: for input Z and target Y, learning $P(X)$ helps learning $P(Y|X)$ because features of X help explain Y

4. *shared factors across tasks*: like previous point, but also works for different Ys

5. *manifolds*: probability mass concentrates in regions with much smaller dimensionality than data itself

6. *natural clustering*: different values of categorical variables are associated with separate manifolds.

7. *temporal and spatial coherence*: consecutive or spatially nearby observations tend to be associated with the same value of relevant categorical concepts or result in small surface move on the surface of the manifold

8. *sparsity*: could mean often many features are 0. could also be that the features are insensitive to small changes in x

9. *simplicity of factor dependencies*: ideally factors are related to each other linearly, or otherwise simply

The process of extracting representations from observations, or inferring latent variables in a probabilistic view of a dataset, is often called **inference**. There are **generative** and **discriminative** models.

Generative models learn representations by modelling the data distribution $p(\boldsymbol{x})$. Such a model can generate realistic examples. Evaluating the conditional distribution $p(y|\boldsymbol{x})$ it done via Bayes rule.

Discriminative models model the conditional distribution $p(y|\boldsymbol{x})$ directly. Discriminative modelling consists of first the inference that extracts latent variables $p(\boldsymbol{v}|\boldsymbol{x})$ which are then used to make downstream decision from those variables $p(y|\boldsymbol{v})$.

The benefit of discriminative models is that the expensive process of learning $p(\boldsymbol{x})$ is avoided. That however makes them harder to evaluate. This is especially evident if you just want a lower dimensional distribution. In the context of reinforcement learning, the model-based approach benefits from generative models as they can be used to generate predictions which can then be used for planning. In the model-free approach, both discriminative and generative models may be used as predictions are not used.

3.1.1 Generative models

3.1.1.1 Probabilistic models

From the probabilistic modeling perspective, feature learning can be interpreted as an attempt to recover a parsimonious set of latent random variables that describe a distribution over the observed data. $p(x, h)$ is the probabilistic model over the joint space of latent variables h and observed data x. Feature values are then the result of an inference process to determine the probability distribution of the latent variables given the data, i.e. $p(h|x)$, a.k.a posterior probability. Learning is the finding the parameters that (locally) maximize the regularized likelihood of the training data.

3.1.1.2 Directed graphical models

Directed latent factor models separately parametrize $p(x|h)$ and the prior $p(h)$ to construct $p(x, h) = p(x|h)p(h)$. They can explain away: a priori independent causes of an event can become nonindependent given the observation of the event. Can conceive them as cause models, where h activations cause the observed x, making h nonindependent. This makes recovering the posterior $p(h|x)$ intractable.

3.1.1.3 Directly learning a parametric map from input to representation

The posterior distribution becomes complicated quickly. Thus approximate inference becomes necessary, which is not ideal. Also, depending on the problem, one needs to derive feature vectors from the distribution. If we want deterministic feature values in the end, we might as well go ahead and use a nonprobabilistic feature learning paradigm. Doing so is particularly desirable for representations for model-free reinforcement learning algorithms: since the data distribution is not explicitly used to make plans, the stochasticity inherent in statistical modelling hinders the ability of the reinforcement learning algorithm to use those representations.

3.1.2 Discriminative models

In discriminative modelling the data distribution is not directly represented. Instead, it is implicit in the representation space. One way to learn discriminative models is through contrastive representation learning. Intuitively, it's learning by comparing. So instead of needing data labels y for datapoints $\boldsymbol{x}$, you need to define a similarity distribution which allows you to sample a positive input $\boldsymbol{x}^+ \sim p^+(\cdot|\boldsymbol{x})$ and a data distribution for a negative input $\boldsymbol{x}^- \sim p^-(\cdot|\boldsymbol{x})$, with respect to an input sample $\boldsymbol{x}$. "Similar" inputs should be mapped close together, and "dissimilar" samples should be mapped further away in the embedding space.

Let's explain how this would work with the example of image-based instance discrimination. The goal is to learn a representation by maximizing agreement of the encoded features (embeddings) between two differently augmented views of the same images, while simultaneously minimizing the agreement between different images. To avoid model maximizing agreement through low-level visual cues, views from the same image are generated through a series of strong image augmentation methods. Let $\mathcal{T}$ be a set of image transformation operations where $t, t' \sim \mathcal{T}$ are two different

transformations sampled independently from $\mathcal{T}$. There transformations include for example cropping, resizing, blurring, color distortion or perspective distortion and their combinations. A $(\boldsymbol{x}_q, \boldsymbol{x}_k)$ pair of query and key views is positive when these 2 views are created with different transformations on the same image, i.e. $\boldsymbol{x}_q = t(\boldsymbol{x})$ and $\boldsymbol{x}_k = t'(\boldsymbol{x})$, and is negative otherwise. A feature encoder $e(\cdot)$ then extracts feature vectors from all augmented data samples $\boldsymbol{v} = e(\boldsymbol{x})$. This is usually ResNet, in which case $\boldsymbol{v} \in \mathcal{R}^d$ is the output of the average pooling layer. Each v is then fed into a projection head $h(\cdot)$ made up of a small multi-layer perceptron to obtain a metric embedding $\boldsymbol{z} = h(\boldsymbol{v})$, where $\boldsymbol{z} \in \mathcal{R}^{d'}$ with $d' < d$. All vectors are then normalized to be unit vectors. Then you take a batch of these metric embedding pairs $\{(\boldsymbol{z}_i, \boldsymbol{z}_i')\}$, with $(\boldsymbol{z}_i, \boldsymbol{z}_i')$ being the metric embeddings of $(\boldsymbol{x}_q, \boldsymbol{x}_k)$ of the same image are fed into the contrastive loss function which does what we said 3 times already. The general form of popular loss function such as InfoNCE and NT-Xent is:

$$\mathcal{L}_i = -\log \frac{\exp(\boldsymbol{z}_i^T \boldsymbol{z}_i'/\tau)}{\sum_{j=0}^{K} \exp(\boldsymbol{z}_i \cdot \boldsymbol{z}_j')/\tau} \tag{3.1}$$

where τ is the temperature parameter. The sum is over one positive and K negative pairs in the same minibatch.

3.1.3 Common representation learning approaches

3.1.3.1 Deterministic autoencoders

Deterministic autoencoders are generative models.

$$h^{(t)} = f_\theta(x^{(t)}) \tag{3.2}$$

There's also the reconstruction $r = g_\theta(h)$, used for the reconstruction error $L(x, r)$ over the training examples. Autoencoder training boils down to finding θ which minimizes:

$$\mathcal{J}_{AE}(\theta) = \sum_t L(x^{(t)}, g_\theta(f_\theta(x^{(t)}))) \tag{3.3}$$

One can tie the weights between the encoder and the decoder (i.e. make the same ones, just reversed).

3.1.3.2 Variational autoencoders

Variational autoencoders marry graphical models and deep learning. The generative model is a Bayesian network of form $p(\boldsymbol{x}|\boldsymbol{z})p(\boldsymbol{z})$, or in the case of multiple stochastic layers, a hierarchy such as: $p(\boldsymbol{x}|\boldsymbol{z}_L)p(\boldsymbol{z}_L|\boldsymbol{z}_{L-1})\cdots p(\boldsymbol{z}_1|\boldsymbol{z}_0)$. Similarly, the recognition model is also a conditional Bayesian network of form $p(\boldsymbol{z}|\boldsymbol{x})$ which can also be a hierarchy of stochastic layers. Inside each conditional may be a deep neural network, e.g. $\boldsymbol{z}|\boldsymbol{x} \sim f(\boldsymbol{x}, \boldsymbol{\epsilon})$ with f being the neural network mapping and $\boldsymbol{\epsilon}$ a noise random variable. Its learning algorithm is a mix of classical (amortized, variational) expectation maximization, but with the reparametrization trick ends up backpropagating through the many layers of the deep neural networks embedded inside it.

We can parametrize conditional distributions with neural networks. VAEs in particular work with *directed* probabilistic models, also know as *probabilistic graphical models* (PGMs) or *Bayesian networks*. The joint distribution over the variables of such models factorizes as a product of prior and conditional distributions:

$$p_{\boldsymbol{\theta}}(\boldsymbol{x}_1, \ldots, \boldsymbol{x}_M) = \prod_{j=1}^{M} p_{\boldsymbol{\theta}}(\boldsymbol{x}_j | Pa(\boldsymbol{x}_j)) \tag{3.4}$$

where $Pa(\boldsymbol{x}_j)$ is the set of parent variables of node j in the directed graph. For root nodes the parents are an empty set, i.e. that distribution is unconditional. Before you'd parametrize each conditional distribution with ex. a linear model, and now we do it with neural networks:

$$\boldsymbol{\eta} = \text{NeuralNet}(Pa(\boldsymbol{x})) \tag{3.5}$$
$$p_{\boldsymbol{\theta}}(\boldsymbol{x} | Pa(\boldsymbol{x})) = p_{\boldsymbol{\theta}}(\boldsymbol{x} | \boldsymbol{\eta}) \tag{3.6}$$

To solve intractabilities, we introduce a parametric *inference model* $q_{\boldsymbol{\phi}}(\boldsymbol{z}|\boldsymbol{x})$. This model is called the *encoder* or *recognition model*/ $\boldsymbol{\phi}$ are called the *variational parameters*. They are optimized s.t.:

$$q_{\boldsymbol{\phi}}(\boldsymbol{z}|\boldsymbol{x}) \approx p_{\boldsymbol{\theta}}(\boldsymbol{z}|\boldsymbol{x}) \tag{3.7}$$

Like a DLVM, the inference model can be almost any directed graphical model:

$$q_{\boldsymbol{\phi}}(\boldsymbol{z}|\boldsymbol{x}) = q_{\boldsymbol{\phi}}(\boldsymbol{z}_1, \ldots, \boldsymbol{z}_M | \boldsymbol{x}) = \prod_{j=1}^{M} q_{bm\phi(\boldsymbol{z}_j | Pa(\boldsymbol{z}_j), \boldsymbol{x})} \tag{3.8}$$

This can also be a neural network. In this case, parameters $\boldsymbol{\phi}$ include the weights and biases, ex.

$$(\boldsymbol{\mu}, \log \boldsymbol{\sigma}) = \text{EncoderNeuralNet}_{\boldsymbol{\phi}}(\boldsymbol{x}) \tag{3.9}$$
$$q_{\boldsymbol{\phi}}(\boldsymbol{z}|\boldsymbol{x}) = \mathcal{N}(\boldsymbol{z}; \boldsymbol{\mu}, \text{diag}(\boldsymbol{\sigma})) \tag{3.10}$$

Typically, one encoder is used to perform posterior inference over all of the datapoints in the dataset. The strategy used in VAEs of sharing variational parameters across datapoints is also called *amortized variational inference*.

3.1.3.3 Deterministic autoencoder regularization

Autoencoders may be employed not only just to learn representations, but to perform additional auxiliary tasks. One such task is denoising: provided a noisified input at the encoder, the decoder outputs a denoisified image as output. Importantly, while this training process results in a denoising autoencoder, it also regularizes the autoencoder. Regularization not only helps with preventing overfitting, but also produces better representations as it encourages smoothness and spatial coherence of when learning. The same result can be accomplished by other data augmentation techniques like random cropping.

Learning VAEs from data poses unanswered theoretical questions and considerable practical challenges. This work proposes a generative model that is simpler, deterministic, easier to train, while retaining some VAE advantages. Namely, the observation is that sampling a stochastic encoder in Gaussian VAE can be interpreted as injecting noise into the input of a deterministic decoder.

The encoder deterministically maps a data point $\boldsymbol{x}$ to the mean $\mu_\phi(\boldsymbol{x})$ and variance $\sigma_\phi(\boldsymbol{x})$ in the latent space. The input to D_θ is then the mean $\mu_\phi(\boldsymbol{x})$ augmented with Gaussian noise scaled by $\sigma_\phi(\boldsymbol{x})$ via the reparametrizing trick. Authors argue that this noise injection is a key factor in having a regularized decoder (noise injection as a mean to regularize neural networks is a well-known technique). Thus training the RAE involves minimizing the simplified loss:

$$\mathcal{L}_{\mathrm{RAE}} = \mathcal{L}_{\mathrm{REC}} + \beta\mathcal{L}_{\boldsymbol{z}}^{\mathrm{RAE}} + \lambda\mathcal{L}_{\mathrm{REG}} \tag{3.11}$$

where $\mathcal{L}_{\mathrm{REG}}$ represents the explicit regularizer for D_θ, and $\mathcal{L}_{\boldsymbol{z}}^{\mathrm{RAE}} = \frac{1}{2}||\boldsymbol{z}||_2^2$, which is equivalent to constraining the size of the learned latent space, which is needed to prevent unbounded optimization. One option for $\mathcal{L}_{\mathrm{REG}}$ is Tikhonov regularization since it is known to be related to the addition of low-magnitude input noise. In this framework this equates to $\mathcal{L}_{\mathrm{REG}} = \mathcal{L}_{L_2} = ||\theta||_2^2$. There's also the **gradient penalty** and **spectral normalization**.

3.2 Representation models for control

In state representation learning the learned features are of low dimension, evolve through time and are depended on actions of an agent. The last point is particularly important because in reinforcement learning, features that do not influence the agent and that can not be influenced by the agent are not relevant for the problem of optimally controlling the agent. Also, simply reducing the dimensionality of the input to a reinforcement learning agent results in a computationally easier learning problem, which can make a difference between the solution being feasible or infeasible. Ideally, state representation learning should be done in an without explicit supervision as it can then be done in tandem with the likewise unsupervised reinforcement learning.

While we assume that state-transitions have the Markov property, partial observability denies the possibility of having a one-to-one correspondence between each observation and state — an object whose position is required may be occluded by another. Thus prior observations have affect the mapping to the current state. Images in particular also do not encode kinematic or dynamic information: to get that crucial information a sequence of images is required. Hence we define the SRL task as learning a representation $\tilde{\boldsymbol{s}}_t \in \tilde{\mathcal{S}}$ of dimension K with characteristics similar to those of true states $\boldsymbol{s}_t \in \mathcal{S}$. In particular, the representation is a mapping of the history of observation to the current state: $\tilde{\boldsymbol{s}}_t = \phi(\boldsymbol{o}_{1:t}$. Actions $\boldsymbol{a}_{1:t}$ and rewards $r_{1:t}$ can also be added to the parameters of ϕ. This can help in extracting only the information relevant for the agent and its task. Often the representation is learned by using the reconstruction loss; $\hat{\boldsymbol{o}}_t$ denotes the reconstruction of $\boldsymbol{o}_t$.

In the context of reinforcement learning, state representations should ideally have the following properties:

- have the Markov property
- be able to represent the current state well enough for policy improvement
- be able to generalize to unseen states with similar features
- be low dimensional

We now discuss different types of models and learning strategies which can be used to learn state representations.

One way to do this is to explicitly use such methods to learn a function which maps from observations to states and then use reinforcement learning methods these learned state representations. This approach is explored in this section, mainly with the help of the [Les+18] overview paper. In this section state representation learning for control in general is discussed. as this will allow for a broader contextualization of our own work.

With representation in learning introduced in general, we can now introduce four different strategies for learning latent space models for control: the autoencoder, the forward model, the inverse model and the model with prior. These models refer to portions of the control problem they are modelling.[1] They can be both discriminative and generative models. In the figures below, the white nodes are inputs and the gray nodes are outputs. The dashed rectangles are fitted around variables with which the loss is calculated.

3.2.1 Autoencoder

The idea behind the autoencoder is to just learn a lower-dimensional embedding of the observation space. This should make the learning problem easier due to the dimensionality reduction. The auto-encoder may be trained to denoise the observations by passing an observation with artificially added noise to the encoder, but then calculating the reconstruction loss on the image without the added noise. Formally this can be written as

$$\boldsymbol{s}_t = \phi(\boldsymbol{o}_t; \theta_\phi) \tag{3.12}$$

$$\hat{\boldsymbol{o}}_t = \phi^{-1}(\boldsymbol{s}_t; \theta_{\phi^{-1}}) \tag{3.13}$$

where θ_ϕ and $\theta_{\phi^{-1}}$ are the parameters learned for the encoder and decoder respectively.

3.2.2 Forward model

The auto-encoder does not encode dynamic information. Since that information is necessary for control, usually a few consecutive observations (or their embeddings) are stacked and passed to the reinforcement learning algorithm. This way the information about the dynamics is implicitly provided. While doing so works, it could be made more efficient by embedding the dynamic information as well. One way to achieve this is to trained a model to predict future state representations. A model can also be observations directly, of course provided that the network in question has a bottleneck layer from which the learned representations can be extracted. Since

---

[1]The term autoencoder is overloaded in this case.

Figure 3.1: Auto-encoder: learned by reconstructing the observation (one-to-one). The observation is the input and the computed state is the vector at the auto-encoder's bottleneck layer, i.e. is the output of the encoder part of the auto-encoder network. The loss is calculated between the true observation and the reconstructing observation (which is obtained by passing the observation though both the encoder and the decoder).

learning on sequential information is difficult and would also benefit from lowering the dimensionality, learning a forward model can be done in two steps: first, learning an auto-encoder to embed individual frames and then learning a predictive model in the embedded space. In the schematic we show the case where predictions are learned from embeddings because it is the structurally more complex scheme. Formally, we have

$$\hat{\tilde{s}}_{t+1} = f(\tilde{s}_t, a_t; \theta_{\mathrm{forward}}) \tag{3.14}$$

The forward model can be constrained to have linear transition between $\tilde{s}_t$ and $\tilde{s}_{t+1}$, thereby imposing simple linear dynamics in the learned state space. Depending on the problem, if this is done well enough, learning a control law can be avoided and instead schemes like model-predictive control can be employed.

Figure 3.2: Forward model: predicting the future state from the state-action pair. The loss is computer from comparing the predicted state against the true next state (the states being the learned states). This can also be done directly by predicting the next observation and comparing against it.

3.2.3 Inverse model

The introducing predictions solves the problem of not embedding the dynamic information. However, not all information in the observation is relevant for control. Consider a computer game where images feature decorative backgrounds — those decorations are irrelevant for playing the game well. If the reconstruction loss is

computed from entire observation, that information is also carried over into the embedded space. However, if the model is trained to predict actions, it is only incentivised to use information which the agent can affect. Thus, due to less information being required, the inverse model should produce a more compact embedding. Formally, we can write this as:

$$\hat{\boldsymbol{a}}_t = g(\tilde{\boldsymbol{s}}_t, \tilde{\boldsymbol{s}}_{t+1}; \theta_{\text{inverse}}) \tag{3.15}$$

If the inverse model is neural network, we can recover the embedding by discarding the last few layers and use their outputs to produce the embeddings.


Figure 3.3: Inverse model: predicting the action between two consecutive states. The loss is computer from comparing the predicted action between two consecutive states against the true action that was taken by the agent between those two states. (the states being the learned states).

3.2.4 Using prior knowledge to constrain the state space

Of course, not everything need be learned in every problem. While in general hand-engineered features are worse than learned ones, there are other ways to provide prior knowledge to the learning system. For example, convolutional neural network by their architecture encode the fact that nearby pixels are related. In the SRL context we already mention the possibility of constraining the model to linear transitions, but there are other available techniques like for example constraining temporal continuity or the principle of causality. Furthermore, priors can be defined as additional objectives or loss functions. For example, additional loss can be provided if embeddings from consequtive observation are drastically different. This is called the slowness principle.

3.2.5 Using hybring objectives

The approaches outlined thus far can be combined into hybrid approaches, for example [Wat+15].

3.3 Model-based reinforcement learning

Like the name suggest, in model-based reinforcement learning a "world model" is learned. While there exists a whole spectrum of methods between pure model-free and model-based ones, the key distinguishing feature of model-based methods is that the learned model is used to *plan* actions. In this case, the task of reinforcement learning in the narrow sense is to learn the values of different states. This then enables calculation of trajectories toward states with high rewards. In model-free methods, only the following action is selected at on iteration of the process because only the transition reward is learned and states these transitions lead to are unknown (not explicitly modelled).

4

Related Work

As said in the introduction, the goal of the thesis is to use state representation learning to increase the efficiency and finals results of model-free reinforcement learning. We are now ready to discuss the specifics of our approach. Firstly, we limit ourselves to image observations and discrete action spaces. In particular, we limit ourselves to Atari57 games as they are common benchmarks in the field for discrete action spaces. As shall be seen in the following text, a lot of recent work in state-representation learning for model-free reinforcement learning has been done in robotics problems with continous action spaces, for example [Yar+19]. Importantly, since we are concerned with finding ways to make reinforcement learning more sample-efficient, we will be using only off-policy algorithms.

Secondly, we are particularly interested in the problem of simultaneous training of the state representations and the policy. The reason for this is that two-step training is often not available because not all state transitions can be observed beforehand. This state of affairs is the natural setting for problems where reinforcement learning is a good solution: the problems where exploration is necessary due to either the high complexity of the dynamics or unanticipatable events. Parallel training of the state representations and the policy necessitates instability in policy training due to the fact the state estimations change even for same observations as the state representation are learned. Hence, related work that focuses on solving or at least ameliorating this issue is of particular importance to our work.

Finally, we want our method to be robust not just in the sense that it works across a wide array of problems, but in the sense that it can be easily added to a variety of reinforcement learning algorithms to a positive effect. In other words, it should function as a module which can be easily added to new algorithms. Furthermore, it should work well with other improvements as those suggested in some of the following related work. To set the context, we begin with by discussing prior work in the Atari environment.

4.1 Reinforcement learning on Atari

Started with [Mni+13]. We already discussed [Hes+18]. Agent 57 [Bad+20] was the first deep RL agent that out performs the standard human benchmark on all 57 Atari games. It was built on on top of the Never Give Up (NGU) agent which utilizes a model-based approach. It combines two ideas: first, the curiosity-driven exploration, and second, distributed deep RL agents, in particular R2D2. The agent was able to balance the learning of different skills that are required to perform well

on such diverse set of games: exploration and exploitation and long-term credit assignment. In order to achieve this a neural network was trained to parameterize a family of policies ranging from very exploratory to purely exploitative, by using adaptive mechanism polices were prioritized throughout the training process.

However, if we convert simulated time to real time, these algorithms can take up to 16000 hours to reach their final performance. Since the goal is not really to solve Atari games, but to find useful general purpose algorithms, the work is still ongoing. The new proposed benchmark is Atari100K: solving the games with with only 100000 transitions. [1] This equates to 2.5 hours of real time.

4.2 Efforts in increasing efficiency in Atari

At the moment of writing, to the authors knowledge, the most efficient algorithm is [Ye+21] which is based on MuZero [Sch+20a] and is a model-based algorithm. However, the title of the most efficient algorithm often switches between a model-based algorithm, a model-free algorithm with state representation learning or similar approaches. We will not discuss model-based approaches, but will discuss some alternative ones as their techniques illuminate the problem.

In particular, this concerns using data-augmentation as a means to directly regularize reinforcement learning. This was first employed in [Las+20] and expanded in [KYF20] and [Yar+21]. In [Las+20], the observations are augmented before they are passed to the policy networks. As we discussed in 3.1.3.3, data-augmentation or noisifying input data functions as strong regularization to feature extractors. The same applies to feature extraction trained just from reinforcement learning. In [KYF20], the same observation is copied and augmented several times. All of these augmented version of the same image are passed through the policy network. The results are then averaged and provide a better estimates than those obtained by a single pass of either non-modified or augmented observation. Thus we may conclude that data augmentation provides benefits to both representation and reinforcement learning.

We now turn to discussing works which utilize unsupervised state representation learning to increase reinforcement learning efficiency.

4.3 State representation learning for efficient model-free learning

Auxiliary losses may be used in a myriad of different ways to help reinforcement learning. In for example [She+16], [Jad+16] or [Pat+17] the same models used for state representations as used to help guide exploration. When, for example, a trained forward predictive model incurs large error, it is reasonable to assume that this happened because a novel state has been encountered. This means that the loss

---

[1]This equates to 400000 frames because the standard is to repeat each action 4 times: this makes learning easier, but also makes sense because humans do not need such small reaction time to solve the games.

can be interpreted as "intrinsic reward" and be added to "extrinsic reward" provided by the environment, yielding an algorithm which encourages exploration.

Of interest to us is the use of auxiliary losses for state representation learning. The specific loss and how it's used depends on the chosen state representation model. In the following subsections some common approaches will be explored.

4.3.1 Deterministic generative models

Perhaps the simplest model to be used for state representation learning on images is an autoencoder trained on reconstruction loss. Using an autoencoder ensures spatial coherence. This idea has been introduced in [LR10]. It did not get traction in reinforcement learning more broadly due to the fact that when the autoencoder is updated, the state representation changes. Unlike regularizing noise which reduces overfitting and incentivizes learning of desirable properties, this noise is destructive. It hinders the ability of the reinforcement learning algorithm to associate states with their values due to the fact that what it is given different numbers as the same state through the course of autoencoder training. To solve this problem, regularization needs to used. In [Yar+19], this was solved by employing the regularizations introduced in [Gho+19], which were already discussed in 3.1.3.3.

A mayor flaw of this approach is the fact that reconstructive loss incentivizes reconstruction of the entire image which contains information irrelevant to the agent. This pertains backgrounds and other object which do not effect state transitions. This does not mean that the obtained representations are not better than raw images, but that they could be made better. Knowing this, we still opted for this approach due to its simplicity and easy of debugging.

4.3.2 Stochastic generative models

Stochastic models can be used to generate predictions which can be used to plan and thus be used in model-based reinforcement learning. However, this does not mean that they can not be used to bolster model-free learning. As discussed in 2.1.2, the formal setting for reinforcement learning is the Markov decision process which is stochastic. Of course, the degree of stochasticity depends on the problem at hand, but given even in a fully deterministic setting stochastic models can be used to deal with epistemic uncertainty. This is further exacerbated in case of partial observability. In [Lee+20], (approximate) variational inference is used to formulate the entire algorithm objective. First, control is formulated as an inference problem and is thereby embedded into the MDP graphical model. From this single graphical model of the problem, the variational distribution of action-dependent state transitions can be factorized into a product of recognition, dynamics and policy terms. As with most approaches which employ stochastic generative models, a variational autoencoder is used to represent the latent (representation) space. It should be noted that without this deep integration with the problem, which enables learning state representation and policies under a single objective, the stochasticity of state representations would hurt the performance of the algorithm. A detailed analysis of this issue can be found in [Yar+19].

4.3.3 Discriminative models

Because we are ultimately only interested in state representations, generative models are not required. Thus it is natural to opt for a discriminative model. Discriminative models can be trained in different ways. In [LSA20], contrastive loss is employed. Another common choice, theoretically investigated in [Rak+21], (TODO:check these 2) is used in [Ana+19] and [Maz+20] is to use mutual information. A particularly promising avenue is to learn discriminative representation models though bootstrapping as introduced in [Gri+20]. This has been employed to learn state representations in [Sch+20b], and in [Mer+22a] where the losses have also been used to incentivize exploration.

These approaches ameliorate problems found in approaches discussed so far: that they avoid both the stochasticity of stochastic generative models and the unnecessary features picked up through reconstruction loss. Learning state representations through bootstrapping is particularly interesting because it is rather flexible with its formulation. In both papers mentioned, the bootstrapping happens through self-predictive loss and is aided with inverse dynamics loss. It would be interesting to integrate this more deeply with an appropriate reinforcement learning algorithm, akin to how stochastic generative models are integrated in the MDP in [Lee+20].

5
Methods

5.1 Formulating our hypotheses

Early attempts to integrate unsupervised state representation learning with reinforcement learning did not yield the expected results. One of the reasons for this was due to the fact that joint training of both algorithms did not work due to the instability caused by different training objectives. Before discussing how to solve this issue, we first need to elaborate why joint training is beneficial. In most recent work, convolutional neural networks are used to extract features from images. This includes reinforcement learning on images: convolutional layers precede linear layers. The convolutional layers, with possible addition of some of the linear layers, constitute the "feature extraction" portion of the policy networks which learn from images. In all of the related work which utilizes state representation learning and is discussed above unsupervised learning is used train the feature extractor portion of the reinforcement learning network. Because unsupervised learning tasks are not trained to extract states in a supervised manner (because the states are inaccessible), the features they learn will not exactly correspond to true states nor will they be able to exactly learn the state transition dynamics. This of course does not mean that approximate state representation can not be sufficiently accurate however. For example, in [Sto+21] the feature extractor trained with augmented temporal contrast (a novel unsupervised learning task introduce in the work) was, on some problems, able to learn features which resulted in sample-efficiency close to training on true states.

On the other hand, by training the policy network to maximize reward, the feature extraction portion of the network will implicitly learn to extract features which are good enough to choose reward-maximizing actions. Thus, logically speaking, the feature extractor trained with just the reinforcement learning signal has to able to extract stateful information for the images, otherwise the agent would not perform well. However, because the reward signal is sparse and contains only indirect information about state, the feature extractor, which contains most of the policy network parameters in simple problems, learns slowly.

This leads us to our first hypothesis 1. Assuming that the learned representations correspond to the images well, thus also meaning that they do not destroy stateful information, then they could be interpreted as partial extractions of states from images. Reinforcement learning feature extraction should then learn more efficiently on top of those representations due to fact that the search space has been constricted. We illustrate this assumption in 5.1. Given the fact that neural networks

have to be overparametrized in order to learn using stochastic gradient descent and their generalization capabilities, the feature extractors for the unsupervised learning and reinforcement learning tasks can be the same network. Furthermore, owing to learning properties of stochastic gradient descent, networks need to be updated with samples taken from all available information. In other words, they "forget" information they have learned unless it continues to be provided in subsequent gradient updates. For this reason we suspect that continuing with unsupervised learning updates even after the unsupervised learning loss initially diminishes will be necessary to ensure the constriction described above. Following this results in joint training of state representation learning and reinforcement learning, which is why we hypothesize that it will yield higher sample-efficiency than other training approaches.


Figure 5.1: Schematic of the feature extractor neural network parameter space. Since unsupervised learning converges much faster, it constricts the search space for the features extracted through reinforcement learning. This constriction is enforced joint training of both unsupervised state representation and reinforcement learning.

Because the objectives are different and because unsupervised learning is much faster, we hypothesize that regularization is necessary 3. Additionally, we hypothesize that representations which better represent true states and state transition dynamics, i.e. the underlying Markov decision process will better condition the reinforcement learning feature extraction task 2.

5.2 Our approach

As stated previously, the goal is to learn effective state representations while training the policy. We opt for a deterministic generative model to learn state representations with, specifically a deterministic autoencoder trained with image reconstruction loss. This is done because performance of a generative model is easier to analyse than discriminative ones. The reinforcement learning algorithm shares its encoder with the encoder for state representation learning (they are the same network, with different heads attached). Our main inspiration for this foundational approach comes from [Yar+19]. However, since we elect to work in discrete action environments,

we choose [Hes+18] as the underlying reinforcement learning algorithm. Since we are not trying to achieve state-of-the-art results, we could have elected another algorithm, but this allowed us to more easily compare our results to those applicable in 4, and also to have faster training times which made experimentation easier. We introduce changes informed by 4 in order to test our hypothesis. Specific papers will be referenced when appropriate. In particular this equates to the following:

1. We test hypothesis 1 with four different training modes.

   (a) Only the reinforcement learning algorithm is trained. This serves as the control case.

   (b) The autoencoder is pretrained, fixed in place, and reinforcement learning algorithm is trained on top of the unsupervised learning features without the ability to updated the encoder. This roughly tells us the quality of features (state representations) obtained through unsupervised learning alone.

   (c) Same as before, but now we update the encoder with reinforcement learning loss. This serves as the control case for the following mode.

   (d) We jointly train the encoder both with unsupervised learning and with reinforcement learning loss. The encoder continues to be updated even after the unsupervised learning loss diminishes.

2. We test hypothesis 2 with the following additional losses.

   (a) Instead of doing unsupervised learning to reconstruct the given observations, we also pass it action and predict the following observations. Since we can not pass the actions before they have been selected by the policy network, they are passed in the decoder (the generative part of the autoencoder). This thus constitutes forward prediction in pixel space.

   (b) Instead of performing forward prediction in pixel space, we now do it in latent space. The autoencoder is trained with reconstructive loss as before, but we introduce an additional network which predicts the latent observation. Thus this loss is calculated separately and it updates the encoder along the other two losses. This serves as a potentially superior version of the previous task as predictions in latent space should be easier to learn that predictions in pixel space.

   (c) Finally, we add inverse dynamics loss on top of the reconstruction, forward predictive and reinforcement learning loss. As shown in [Pat+17], this should encourage the state representation to focus the parts of the observation related to the agent and thus further condition state representations. As argued in [Rak+21], this loss alone can not be represent the MDP so it is not considered in isolation.

3. We test hypothesis 3 with the following regularization techniques.

   (a) Following the analysis carrier out in [Yar+19], we begin by using regularization techniques introduced in [Gho+19]. Since we found that joint training of state representations and reinforcement learning does not work without at least L2 latent space loss, we use it all joint training tests.

   (b) Informed by regularization effects of denoising autoencoders, and by great success achieved with the random shift augmentation in [KYF20; Yar+21] and other recent work, we too use random shift augmentation on obser-

vations in unsupervised learning updates. This serves as our test.

5.3   Environment and Preprocessing

We perform a comprehensive evaluation of our proposed method on the Arcade Learning Environment [Bel+13b], which is composed of 57 Atari games. The challenge is to deploy a single algorithm and architecture, with a fixed set of hyperparameters, to learn to play all the games given embedded latent space representation of the environment from auto encoder and game rewards. This environment is very demanding because it is both comprised of a large number of highly diverse games and the observations are high-dimensional.

Working with raw Atari frames, which are 210 x 160 pixel pictures with a 128 color palette, is computationally expensive, therefore we do a basic preprocessing step to reduce the input dimensionality. The raw frames are down sampled to a 84 x 84 picture and transforming their RGB representation to gray-scale. This is standard practice in the field and it makes the training faster without fundamentally simplifying the problem.

We constrict ourselves to games in which there is no partial observability and where the reward is relatively dense. This is because we are not focusing on exploration which is necessary to solve some of the games. Methods known to use which formulate intrinsic rewards are compatible with our approach and can be in fact integrated with it like in [Mer+22b].

5.4   Implementation

5.4.1   Tianshou

We implement our approach as a modular extension in [Wen+21]. The purpose of this choice is to make our work easily accessible and extendable. Tianshou is an actively updated reinforcement learning library and offers the highest number reinforcement learning algorithms on the market. It is based on Python and PyTorch. It is well documented, provides tests and type hints. It offers gym wrappers for the most popular reinforcement learning benchmarks, including Atari. It supports vectorization for all environments.

The guiding principle of Tianshou is to abstract the reinforcement learning problem as much as possible because good abstractions are the basis of modularity. This is done by splitting reinforcement learning algorithms into the following modules:

- `Batch`
- `ReplayBuffer`
- `Policy`
- `Collector`
- `Trainer`

These modules interact in the way depicted in **??**.

`Batch` is designed to store and manipulate "hierarchical named tensors". Hierarchical named tensors are a set of tensors whose name forms a hierarchy. In essence,

Figure 5.2: Tianshou concepts

they are nested dictionaries.

There are 7 reserved keys in `Batch`:

1. `obs` — observation at step t
2. `act` — action at step t
3. `rew` — reward at step t
4. `done` — done flag at step t
5. `obs_next` — observation at step $t + 1$
6. `info` — info at step t
7. `policy` — data computed by policy at step t

`ReplayBuffer` stores experience data. It's purpose is to manage `Batch`. All data is stored in a circular queue.

There are different classes of policies, but all policies must inherit from `BasePolicy`. Typical function are the following ones:

1. `__init__()`
2. `forward()` — compute action with given observation
3. `process_fn()` — pre-process data from the replay buffer
4. `learn()` — update policy with a given batch of data
5. `post_process_fn()` — pre-process data from the replay buffer
6. `update()` — this one does it all: samples from buffer, pre-processes data (ex. computing the n-step return), learn from data and post-proces sthe data (ex. update the prioritized replay buffer)

`Collector`'s task is to interact with the environment and store the observed transitions in `ReplayBuffer`.

5.4.2 Trainer

's task is to balance environment interactions by calling on `Collector` and updating the agent by calling on `Policy`.

5.4.3 Implementing state representation learning in Tianshou

Surprisingly, Tianshou lacks state representation learning algorithms which is why believe it could benefit from our work. We implemented state representation learning as a `Policy` which takes a reinforcement learning `Policy` as an argument. This makes our implementation truly modular and we had to change less than 5 lines DQN implementation for it to work. We believe that this was due to a minor design error in n-step return implementation, rather than our lack of foresight. Apart from this we used Tianshou's design to our advantage. By having the state representation learning wrapped around reinforcement learning, we manipulate the batches sent to the reinforcement learning algorithm and thus keep it fully encapsulated. Data augmentation is implemented passing appropriate preprocessing functions to the `Trainer`. The user only needs to ensure matching network dimensions to utilize state representation learning algorithms and pass the appropriate flag to use the one we implemented.

5.5 Hyperparameters

The first step in the process is to collect data to train the auto-encoder. we run a data collection module to generate the 100000 frame for each stated under the result section. The raw images are transformed to tensors and then trained a variational autoencoder with the objective of re-constructing the original image fed to the network. The auto-encoder was trained for maximum 100 epochs. When reconstructing an image with a network bottleneck, the encoder is forced to compress the original image to a smaller dimensional vector in the latent space.
[By compressing the raw pixels environment to smaller dimensional vector in the latent space we aim to improve the training time it takes for the integrated Rl agent developed by [16] and the shift in the latent space representation and its impact on the RL agent learning performance.We show this in more detail in the following sections.]

6
Results

Are being turned into graphs at the moment.

7
Conclusion

Bibliography

[LR10]   Sascha Lange and Martin Riedmiller. "Deep auto-encoder neural networks in reinforcement learning". In: *The 2010 international joint conference on neural networks (IJCNN)*. IEEE. 2010, pp. 1–8.

[Bel+13a] M. G. Bellemare et al. "The Arcade Learning Environment: An Evaluation Platform for General Agents". In: *Journal of Artificial Intelligence Research* 47 (June 2013), pp. 253–279.

[Bel+13b] Marc G Bellemare et al. "The arcade learning environment: An evaluation platform for general agents". In: *Journal of Artificial Intelligence Research* 47 (2013), pp. 253–279.

[Mni+13]   Volodymyr Mnih et al. "Playing Atari with Deep Reinforcement Learning". In: (2013). cite arxiv:1312.5602Comment: NIPS Deep Learning Workshop 2013. URL: http://arxiv.org/abs/1312.5602.

[Mni+15]   Volodymyr Mnih et al. "Human-level control through deep reinforcement learning". In: *Nature* 518.7540 (Feb. 2015), pp. 529–533. ISSN: 00280836. URL: http://dx.doi.org/10.1038/nature14236.

[Wat+15]   Manuel Watter et al. "Embed to control: A locally linear latent dynamics model for control from raw images". In: *Advances in neural information processing systems* 28 (2015).

[Jad+16]   Max Jaderberg et al. "Reinforcement learning with unsupervised auxiliary tasks". In: *arXiv preprint arXiv:1611.05397* (2016).

[She+16]   Evan Shelhamer et al. "Loss is its own reward: Self-supervision for reinforcement learning". In: *arXiv preprint arXiv:1612.07307* (2016).

[Pat+17]   Deepak Pathak et al. "Curiosity-driven exploration by self-supervised prediction". In: *International conference on machine learning*. PMLR. 2017, pp. 2778–2787.

[Hes+18]   Matteo Hessel et al. "Rainbow: Combining improvements in deep reinforcement learning". In: *Thirty-second AAAI conference on artificial intelligence*. 2018.

[Les+18]   Timothée Lesort et al. "State representation learning for control: An overview". In: *Neural Networks* 108 (2018), pp. 379–392.

[SB18]   Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.

[Ana+19]   Ankesh Anand et al. "Unsupervised state representation learning in atari". In: *Advances in Neural Information Processing Systems* 32 (2019).

[Gho+19]   Partha Ghosh et al. "From variational to deterministic autoencoders". In: *arXiv preprint arXiv:1903.12436* (2019).

[Yar+19]   Denis Yarats et al. "Improving sample efficiency in model-free rein-
           forcement learning from images". In: *arXiv preprint arXiv:1910.01741*
           (2019).

[Bad+20]   Adrià Puigdomènech Badia et al. "Agent57: Outperforming the atari
           human benchmark". In: *International Conference on Machine Learn-
           ing.* PMLR. 2020, pp. 507–517.

[Gri+20]   Jean-Bastien Grill et al. "Bootstrap your own latent-a new approach
           to self-supervised learning". In: *Advances in Neural Information Pro-
           cessing Systems* 33 (2020), pp. 21271–21284.

[KYF20]   Ilya Kostrikov, Denis Yarats, and Rob Fergus. "Image augmentation
           is all you need: Regularizing deep reinforcement learning from pixels".
           In: *arXiv preprint arXiv:2004.13649* (2020).

[LSA20]   Michael Laskin, Aravind Srinivas, and Pieter Abbeel. "Curl: Contrastive
           unsupervised representations for reinforcement learning". In: *Interna-
           tional Conference on Machine Learning.* PMLR. 2020, pp. 5639–5650.

[Las+20]   Misha Laskin et al. "Reinforcement learning with augmented data". In:
           Advances in Neural Information Processing Systems 33 (2020), pp. 19884–
           19895.

[Lee+20]   Alex X Lee et al. "Stochastic latent actor-critic: Deep reinforcement
           learning with a latent variable model". In: *Advances in Neural Infor-
           mation Processing Systems* 33 (2020), pp. 741–752.

[Maz+20]   Bogdan Mazoure et al. "Deep reinforcement and infomax learning". In:
           Advances in Neural Information Processing Systems 33 (2020), pp. 3686–
           3698.

[Sch+20a] Julian Schrittwieser et al. "Mastering atari, go, chess and shogi by
           planning with a learned model". In: *Nature* 588.7839 (2020), pp. 604–
           609.

[Sch+20b] Max Schwarzer et al. "Data-efficient reinforcement learning with self-
           predictive representations". In: *arXiv preprint arXiv:2007.05929* (2020).

[Lev21]   Sergey Levine. *Deep Reinforcement Learning course (CS 285) lectures.*
           Fall 2021. URL: https://www.youtube.com/playlist?list=PL_
           iWQOsE6TfXxKgI1GgyV1B_Xa0DxE5eH.

[Rak+21]   Kate Rakelly et al. "Which Mutual-Information Representation Learn-
           ing Objectives are Sufficient for Control?" In: *Advances in Neural In-
           formation Processing Systems* 34 (2021).

[Sto+21]   Adam Stooke et al. "Decoupling representation learning from reinforce-
           ment learning". In: *International Conference on Machine Learning.*
           PMLR. 2021, pp. 9870–9879.

[Wen+21]   Jiayi Weng et al. "Tianshou: A Highly Modularized Deep Reinforce-
           ment Learning Library". In: *arXiv preprint arXiv:2107.14171* (2021).

[Yar+21]   Denis Yarats et al. "Mastering visual continuous control: Improved
           data-augmented reinforcement learning". In: *arXiv preprint arXiv:2107.09645*
           (2021).

[Ye+21]   Weirui Ye et al. "Mastering atari games with limited data". In: *Ad-
           vances in Neural Information Processing Systems* 34 (2021).

[Mer+22a]   Astrid Merckling et al. "Exploratory State Representation Learning".
            In: *Frontiers in Robotics and AI* 9 (2022).
[Mer+22b]   Astrid Merckling et al. "Exploratory State Representation Learning".
            In: *Frontiers in Robotics and AI* 9 (2022). ISSN: 2296-9144. DOI: 10.
            3389/frobt.2022.762051. URL: https://www.frontiersin.org/
            article/10.3389/frobt.2022.762051.

Bibliography

A
Appendix 1

(UNFINISED)
//UPDATE THE FORMAT LATER

1. Agent: It is an assumed entity which performs actions in an environment to gain some reward.
2. Environment (e): A scenario that an agent has to face.anything the agent cannot change arbitrarily is considered to be part of the environment.
3. Reward (R): An immediate return given to an agent when he or she performs specific action or task.
4. State (s): State refers to the current situation returned by the environment.
5. Policy (π): It is a strategy which applies by the agent to decide the next action based on the current state.
6. Value (V): It is expected long-term return with discount, as compared to the short-term reward.
7. Value Function: It specifies the value of a state that is the total amount of reward. It is an agent which should be expected beginning from that state.
8. Model of the environment: This mimics the behavior of the environment. It helps you to make inferences to be made and also determine how the environment will behave.
9. Model based methods: It is a method for solving reinforcement learning problems which use model-based methods.
10. Q value or action value (Q): Q value is quite similar to value. The only difference between the two is that it takes an additional parameter as a current action.